Hao Zhang

The Robotics Institute
School of Computer Science
Carnegie Mellon University
Pittsburgh, PA 15213

Phone: (412) 499-1581
Email: hao@cs.cmu.edu
sjtu.haozhang@gmail.com

Research Interest

Scalable and distributed machine learning

Machine learning systems

Automated machine learning (AutoML)

Deep learning with structures

Large-scale ML applications in computer vison and natural language processing

Education

Robotics Institute, School of Computer Science, Carnegie Mellon University
Ph.D. in Robotics, 2014 - 2020
M.S. in Robotics, 2014 - 2016
Advisor: Prof. Eric Xing

Department of Computer Science and Engineering, Shanghai Jiao Tong University
M.S. in Computer Science, 2011 - 2014
Advisor: Prof. Liqing Zhang

School of Computer Science and Engineering, South China University of Technology
B.E. in Computer Science, the Elite Class of Computer Science, 2008 - 2011

Working Experience

Petuum Inc, Pittsburgh

Director of Scalable ML, December 2017 – Present.

Research Scientist, May 2017 – Present.

Tech Lead, May 2017 – December 2017.

Consultant, Jul 2016 – May 2017.

Microsoft Research Asia, Beijing

Research Intern, July 2013 - January 2014.

Microsoft, Shanghai

SDE Intern, September 2012 - April 2013.

Teaching

Teaching Assistant in Carnegie Mellon University

10-708, Probabilistic Graphical Models, Spring 2019. Instructor: Prof. Eric Xing.

16-791, Applied Data Science, Spring 2019. Instructor: Prof. Artur Dubrawski.

10-701, Introduction to Machine Learning, Fall 2015. Instructors: Prof. Eric Xing and Prof. Ziv Bar-Joseph.

Teaching Assistant in Shanghai Jiao Tong University

Artificial Intelligence, Spring 2012. Instructor: Prof. Liqing Zhang.

Publications

Refereed Conference Papers

[1] **Hao Zhang**†, Haowen Xu†, Zhiting Hu, Xiaodan Liang, Ruslan Salakhutdinov, and Eric P. Xing. AutoLoss: Learning Discrete Schedules for Alternate Optimization. In *The 7th International Conference on Learning Representations*. (**ICLR 2019**) († indicates equal contributions).

[2] Wei Dai, Yi Zhou, Nanqing Dong, **Hao Zhang**, and Eric P. Xing. Toward Understanding the Impact of Staleness in Distributed Machine Learning. In *The 7th International Conference on Learning Representations*. (**ICLR 2019**).

[3] Xiaodan Liang, Zhiting Hu, **Hao Zhang**, Liang Lin, and Eric P. Xing. Symbolic Graph Reasoning Meets Convolutions. In *The 32nd Conference on Neural Information Processing Systems*. (**NIPS 2018**).

[4] Xiaodan Liang, **Hao Zhang**, and Eric P. Xing. Generative Semantic Manipulation with Contrasting GAN. In *2018 European Conference on Computer Vision*. (**ECCV 2018**).

[5] Wei Dai, Nanqing Dong, Zeya Wang, Xiaodan Liang, **Hao Zhang**, and Eric P. Xing. SCAN: Structure Correcting Adversarial Network for Chest X-rays Organ Segmentation. In *2018 Deep Learning in Medical Image Analysis and Multimodal Learning for Clinical Decision Support*.

[6] **Hao Zhang**†, Shizhen Xu†, Graham Neubig, Qirong Ho, Guangwen Yang, and Eric P. Xing. Cavs: An Efficient Runtime System for Dynamic Neural Networks. In *2018 USENIX Annual Technical Conference*. (**ATC 2018**, **Oral**, **AISys@SOSP 2017**, **MLSys@NIPS 2017**) († indicates equal contributions).

[7] **Hao Zhang**†, Zhijie Deng†, Xiaodan Liang, Luona Yang, Shizhen Xu, Jun Zhu, and Eric P. Xing. Structured Generative Adversarial Networks. In *The 31st Conference on Neural Information Processing Systems*. (**NIPS 2017, NVIDIA Pioneer Research Award**) († indicates equal contributions).

[8] Xiaodan Liang, Zhiting Hu, **Hao Zhang**, Chuang Gan, and Eric P. Xing. Recurrent Topic-Transition GAN for Visual Paragraph Generation. In *2017 International Conference on Computer Vision*. (**ICCV 2017**).

[9] **Hao Zhang**, Zeyu Zheng, Shizhen Xu, Wei Dai, Qirong Ho, Xiaodan Liang, Zhiting Hu, Jinliang Wei, Pengtao Xie, and Eric P. Xing. Poseidon: An Efficient Communication Architecture for Distributed Deep Learning on GPU Clusters. In *2017 USENIX Annual Technical Conference*. (**USENIX ATC 2017**, **Oral**).

[10] **Hao Zhang**, Zhiting Hu, Jinliang Wei, Pengtao Xie, Gunhee Kim, Qirong Ho, and Eric Xing. Poseidon: A System Architecture for Efficient GPU-based Deep Learning on Multiple Machines. In *2016 USENIX Annual Technical Conference*. (**USENIX ATC 2016**, **Poster**, **MLSys@ICML 2016**, **Spotlight**).

[11] **Hao Zhang**, Zhiting Hu, Yuntian Deng, Mrinmaya sachan, Zhicheng Yan, and Eric Xing. Learning Concept Taxonomies from Multi-modal Data. In *The 54th Annual Meeting of the Association for Computational Linguistics*. (**ACL 2016**, **Oral**).

[12] Zhicheng Yan, **Hao Zhang**, Baoyuan Wang, Sylvain Paris, and Yizhou Yu. Automatic Photo Adjustment Using Deep Neural Networks. In *2016 International Conference on Computational Photography*. (**ICCP 2016**, **Invited Poster**).

[13] Henggang Cui, **Hao Zhang**, Gregory R. Ganger, Phillip B. Gibbons, and Eric Xing. GeePS: Scalable Deep Learning on Distributed GPUs with a GPU-specialized Parameter Server. In *2016 European Conference on Computer Systems*. (**EuroSys 2016**).

[14] Jincheng Mei, **Hao Zhang**, and Baoliang Lu. On the Reducibility of Submodular Functions. In *The 19th International Conference on Artificial Intelligence and Statistics*. (**AISTATS 2016**).

[15] Zhicheng Yan, **Hao Zhang**, Robinson Piramuthu, Vignesh Jagadeesh, Dennis DeCoste, Wei Di, and Yizhou Yu. HD-CNN: Hierarchical Deep Convolutional Neural Network for Large Scale Visual Recognition. In *2015 International Conference on Computer Vision*. (**ICCV 2015**).

[16] **Hao Zhang**, Gunhee Kim, and Eric Xing. Dynamic Topic Modeling for Monitoring Market Competition from Online Text and Image Data. In *The 21st ACM SIGKDD Conference on Knowledge Discovery and Data Mining*. (**KDD 2015**, **Oral**).

[17] Ye Liu, **Hao Zhang**, and Liqing Zhang. Common Spatial-Spectral Boosting Pattern for Brain-Computer Interface. In *The 21st European Conference on Artificial Intelligence*. (**ECAI 2014**).

[18] **Hao Zhang**, Ye Liu, Jianyi Liang, Jianting Cao, and Liqing Zhang. Gaussian Mixture Modeling in Stroke Patients' Rehabilitation EEG Data Analysis. In *2013 Annual International Conference of the IEEE Engineering in Medicine and Biology Society*. (**EMBC 2013**).

[19] Ye Liu, Mingfen Li, **Hao Zhang**, Junhua Li, Jie Jia, Yi Wu, Jianting Cao, and Liqing Zhang. Single-Trial Discrimination of EEG Signals for Stroke Patients: A General Multi-Way Analysis. In *2013 Annual International Conference of the IEEE Engineering in Medicine and Biology Society*. (**EMBC 2013**).

Refereed Journal Papers

[1] Zhicheng Yan, **Hao Zhang**, Baoyuan Wang, Sylvain Paris, and Yizhou Yu. Automatic Photo Adjustment Using Deep Neural Networks. In *ACM Transactions on Graphics*. (**TOG Vol 35, 2016**).

[2] Ye Liu, **Hao Zhang**, Min Chen, and Liqing Zhang. A Boosting-Based Spatial-Spectral Model for Stroke Patients' EEG Analysis in Rehabilitation Training. In *IEEE Transactions on Neural System and Rehabilitation Engineering*, 2015.

[3] Ye Liu, Mingfen Li, **Hao Zhang**, Hang Wang, Junhua Li, Jie Jia, Yi Wu, Jianting Cao, and Liqing Zhang. A Tensor-based Scheme for Stroke Patients' Motor Imagery EEG Analysis in BCI-FES Rehabilitation Training. In *Journal of Neuroscience Methods*, 2013.

Preprints

[1] **Hao Zhang**, Christy Li, Zhijie Deng, Xiaodan Liang, Lawrence Carin, and Eric P. Xing. AutoSync: Learning to Synchronize for Data-Parallel Distributed Deep Learning. 2020.

[2] **Hao Zhang**, Christy Li, Zhijie Deng, Aurick Qiao, Qirong Ho, and Eric P. Xing. Oceanus: A Composable and Automated Synchronization System for Distributed Deep Learning. 2020.

[3] Aurick Qiao, Willie Neiswanger, **Hao Zhang**, Qirong Ho, Greg Ganger, and Eric P. Xing. Pollux: Co-adaptive Cluster Scheduling for Goodput-Optimized Deep Learning. 2020.

[4] Hao Wang, Xiaodan Liang, **Hao Zhang**, Dit-Yan Yeung, and Eric P. Xing. ZM-Net: Real-time Zero-shot Image Manipulation Network. In **arXiv 1703.07255, 2017**.

[5] Zhicheng Yan, **Hao Zhang**, Yangqing Jia, Thomas Breuel, and Yizhou Yu. Combining the Best of Convolutional Layers and Recurrent Layers: A Hybrid Network for Semantic Segmentation. In **arXiv 1603.04871, 2016**.

Open Source Machine Learning Software

DyNet: The Dynamic Neural Network Toolkit.
Source code available at https://github.com/clab/dynet/.

GeePS: Scalable Deep Learning on Distributed GPUs with a GPU-specialized Parameter Server.
Source code available at https://github.com/cuihenggang/geeps.

Poseidon: Distributed Deep Learning Framework on Petuum.
Software available at http://poseidon-release.readthedocs.io/en/latest/.

Awards and Honors

Nvidia Pioneer Research Award, 2017.

Excellent Graduates (top 5%), Shanghai Jiao Tong University, 2014.

Scholarship for Graduates, Shanghai Jiao Tong University, 2011 - 2014.

Google Excellence Scholarship (National wide), Google Inc., 2013.

Early Graduate Honor (top 1%), South China University of Technology, 2011.

Excellent Undergraduates, South China University of Technology, 2008 - 2011.

1st Class Scholarship (top 10%), South China University of Technology, 2008 - 2011.

Professional Service

Program Committee. AAAI 2020, UAI 2019, UAI 2018.

Reviewer. ICLR 2020, AISTATS 2020, NIPS 2019, ICML 2019, ICLR 2019, NIPS 2018, NACCL HLT 2018, SOSP 2017, CVPR 2016, ICCV 2015, TPAMI, SCIS, IET Computer Vision, MVAP, TCC, VLDB.

Volunteer. ICML 2016, KDD 2015.

Selected Talks

International Conference on Machine Learning, July 2019. *Arion: a Next-generation Distributed Deep Learning Virtual Machine*.

USENIX Annual Technical Conference, July 2018. *Cavs: A Vertex-centric Programming Interface for Dynamic Neural Networks*.

Machine Learning System Workshop, December 2017. *Cavs: A Vertex-centric Programming Interface for Dynamic Neural Networks*.

USENIX Annual Technical Conference, July 2017. *Poseidon: An Efficient Communication Architecture for Distributed Deep Learning on GPU Clusters*.

Annual meeting of the Association for Computational Linguistics, August 2016. *Learning Concept Taxonomies from Multi-modal Data*.

Database Seminar, Carnegie Mellon University, Feb 2016. *Dynamic Topic Modeling for Monitoring Market Competition from Online Text and Image Data*.

Machine Learning Lunch Seminar, Carnegie Mellon University, Feb 2015. *Poseidon: A System Architecture for Large-scale Distributed Deep Learning on GPU Clusters with Commodity Hardware*.

ACM SIGKDD Conference on Knowledge Discovery and Data Mining, Sydney, Aug 2015. *Dynamic Topic Modeling for Monitoring Market Competition from Online Text and Image Data*.

Patents

A System with Hybrid Communication Strategy for Large-scale Distributed Deep Learning. US Patent US0330276A1, 2018.

Structure Correcting Adversarial Network for Chest X-Rays Organ Segmentation. US Patent US0276825A1, 2018.